

MSDM 6980 Computational Modeling and Simulation Project

Predictive Modeling for Non-Performing Assets: A Comparative Analysis between ML and LLMs

This report presents a data-driven framework for NPA debt collection. By benchmarking ML models against LLMs and integrating economic optimization, it demonstrates that calibrated traditional ML maximizes net recovery and significantly outperforms LLM-based approaches

Student: ZHU TIANYI

Student ID: 21229582

Supervisor: Dr. Hanson WONG

Date: May 2026

1 Introduction

1.1 Background and Motivation

Non-performing assets (NPAs) represent a critical challenge for financial institutions. Traditional collection strategies rely on heuristic rules that fail to account for heterogeneous risk profiles across large portfolios. The core problem is **prediction under extreme class imbalance**: in a typical unsecured debt portfolio, only 5–15% of accounts ultimately repay, yet identifying this minority accurately determines whether recovery operations are profitable.

This study compares five modeling paradigms on a real-world Hong Kong NPA portfolio:

1. **Logistic Regression (LR)** — interpretable linear baseline;
2. **Random Forest (RF)** — non-linear ensemble via decision trees;
3. **XGBoost** — gradient boosting for tabular data;
4. **Multi-Layer Perceptron v2 (MLP v2)** — deep network with residual and feature-interaction layers;
5. **Large Language Models (LLMs)** — 26 frontier configurations across zero-shot, few-shot, and fine-tuned regimes.

Beyond model development, this work integrates Platt Scaling calibration and an economic optimization framework that translates predicted probabilities into tiered collection actions. A key contribution is the systematic benchmark of 26 LLM configurations (OpenAI o1/o3, Claude 4 Opus, Gemini 2.5 Pro, DeepSeek-R1, etc.) against four trained ML models for structured financial prediction.

1.2 Research Objectives

- (O1) Develop and compare four ML models for binary repayment prediction within a 3-year window.
- (O2) Evaluate models by discriminative power (ROC-AUC), calibration (Brier Score), economic outcome (net recovery), and ROI.
- (O3) Design a tiered collection strategy assigning accounts to Agent Call, Auto-Dialer, SMS/Email, or Write-off based on calibrated probabilities and cost constraints.
- (O4) Benchmark frontier LLM approaches (zero-shot reasoning, few-shot, fine-tuned) against traditional ML to assess whether “AI-as-judge” can match specialized models.

1.3 Paper Organization

Section 2 documents the student’s contributions. Section 3 reviews key literature. Section 4 describes data, preprocessing, model architectures, and evaluation. Section 5 presents model comparisons, LLM-vs-ML benchmarks, and queue optimization. Section 6 discusses findings and limitations. Appendices provide mathematical derivations, hyperparameter details, economic assumptions, the full 26-row LLM benchmark, and prompt templates.

2 Student’s Unique Contributions

End-to-End ML Pipeline. Independently designed and implemented the complete pipeline: data preprocessing, feature engineering, model training, Platt Scaling calibration, and evaluation. A key domain-aware decision was treating 1,165 missing values in `last_pay_date_client_closing_m` as “never paid” rather than imputing with mean/median, preserving predictive signal in the missingness pattern.

Frontier LLM-vs-ML Benchmark Framework. Proposed and executed a systematic benchmark comparing **26 frontier LLM configurations** (OpenAI o1/o3, Claude 4 Opus, Gemini 2.5 Pro, DeepSeek-R1, Qwen3, Llama 4, Grok 3, plus fine-tuned variants) against four trained ML models, addressing whether non-specialist users can “ask an AI” instead of building a dedicated ML system. Key findings: best zero-shot reasoning LLM (Claude 4 Opus + thinking, AUC = 0.694) closes most of the discrimination gap; best fine-tuned LLM (FT GPT-4.1-mini, AUC = 0.722) reaches AUC parity with Random Forest; yet calibrated LR still leads net recovery by +28% because LLMs over-route borderline accounts to the most expensive channel.

Domain-Calibrated CoT Prompt. Designed a structured prompt template that injects the 9.15% portfolio base rate, the cost grid (Agent 85, Dialer 12, SMS 1.5 HKD), monotonicity priors on age/balance, and forces JSON output. This prompt change alone moved Claude 3.7 Sonnet ZS from AUC 0.625 to 0.671 and reduced self-reported ROI error from >200% to <15%.

Economic Collection Queue Optimization. Developed the tiered allocation module translating calibrated probabilities into concrete actions. The optimized strategy concentrates 50.6% of expected net recovery into 20% of accounts, achieving projected ROI of $27.2\times$ on HKD 83K contact cost, and identified an inverse balance–repayment relationship (accounts <HKD \$25K achieve >11% repayment; those >\$200K achieve 0%).

Interactive Dashboard. Built a 9-tab HTML/JavaScript dashboard (Chart.js) visualizing all model outputs for non-technical stakeholders.

3 Literature Review

Machine learning for credit risk dates to Altman [1] on discriminant analysis. Hand and Henley [4] showed non-linear classifiers substantially outperform linear models on complex financial data, motivating our multi-model approach.

Key methodological references include: Brown and Mues [2] on class imbalance; Niculescu-Mizil and Caruana [7] demonstrating that boosted trees require post-hoc calibration via Platt Scaling [8]; Verbraken et al. [9] advocating profit-based metrics over accuracy; and Grinsztajn et al. [3] finding tree-based methods remain competitive against deep learning on tabular data.

LLMs for Tabular Data. Kuzmin et al. [5] benchmarked LLMs across 18 tabular datasets and found they underperform GBDTs due to suboptimal numerical tokenization, context window constraints, and lack of tabular inductive bias. Fine-tuned LLMs have narrowed this gap [10]; the 2025 reasoning-tuned generation (o1/o3, Claude 4 Opus, Gemini 2.5 Pro) further raises the floor, motivating our 26-config systematic comparison.

4 Data and Methodology

4.1 Dataset

The analysis uses a proprietary portfolio of **16,048** delinquent unsecured debt accounts from a Hong Kong financial institution. Each record is labeled with a binary target: repayment within three years ($\text{payer_3yr} \in \{Y, N\}$). Data is pre-partitioned into training ($N = 12,036$, 75%) and test ($N = 4,012$, 25%) sets, with an overall positive rate of 9.65%.

Table 1: Input Feature Schema (14 features + 1 target)

Name	Type	Missing	Description
loan_type	cat.	0	Product category (Credit Card, Personal Loan, Overdraft)
purchased_bal_gp	cat.	0	Balance bucket group (8 ordinal levels)
district	cat.	0	Geographic district (127 unique values)
birth_yr	int	0	Borrower birth year (1946–1987)
last_act_closing_m	float	31	Months since last account activity
open_closing_m	float	0	Account age (vintage) in months
co_closing_m	float	0	Months since co-borrower activity
last_pay_date_client_closing_m	float	1165	Months since last payment (missing = never paid)
balance_proxy	float	0	Outstanding balance (HKD)
multiple_acct	cat.	0	Holds multiple accounts? (Y/N)
home_phone_flag	int	0	Has registered home phone? (0/1)
mobile_phone_flag	int	0	Has mobile phone? (0/1)
payer_3yr	bin.	0	Target: repaid within 3 years? (Y/N)

4.2 Preprocessing

- (i) **Missing values:** `last_pay_date_client_closing_m` nulls filled with $\max(+999)$; binary flag `never_paid_to_client_flag` created.
- (ii) **Encoding:** one-hot for LR/MLP; ordinal for tree-based models. `birth_yr` $\rightarrow$ `age`. Numeric features z -score standardized for MLP only.
- (iii) **Split:** training set divided into train (60%), validation (20%), calibration (20%); calibration fold reserved exclusively for Platt Scaling.

4.3 Model Architectures

Logistic Regression. Linear baseline with ℓ_2 regularization ($C = 10.0$), balanced class weights, L-BFGS optimization.

Random Forest. $B = 500$ bootstrap-aggregated trees with $\sqrt{d}$ random features per split; max

depth = 12, min leaf size = 8.

XGBoost. $K = 200$ gradient-boosted trees; learning rate 0.05, max depth 6, subsample 0.85, `scale_pos_weight` ≈ 9.18 .

MLP v2. Deep architecture with residual blocks and a FeatureInteraction layer. Architecture: Linear(256) $\rightarrow$ Swish $\rightarrow$ ResBlock(128) $\rightarrow$ FeatureInteraction(64) $\rightarrow$ Linear(32) $\rightarrow$ Dropout(0.3) $\rightarrow$ Output(1). Trained with AdamW, OneCycleLR, label smoothing ($\alpha = 0.03$), gradient clipping (3.0). Detailed derivation in Appendix B.

4.4 Probability Calibration

Raw model outputs are recalibrated via Platt Scaling [8] using 5-fold out-of-fold predictions on a dedicated calibration set:

$$P_{\text{calibrated}}(y = 1|s) = \frac{1}{1 + \exp(A \cdot s + B)}, \quad (1)$$

where s is the raw score and (A, B) are fit via maximum likelihood. Full derivation in Appendix B.

4.5 Evaluation Metrics

Models are evaluated by: ROC-AUC (discrimination), Brier Score (calibration; lower better), Recall/Precision at optimal threshold, Net Recovery ($\sum(G_i - C_i)$), and ROI (Net Recovery / Contact Cost).

4.6 Collection Queue Optimization

Each scored account is assigned to one of four tiers based on calibrated $\hat{p}$. Thresholds and capacity constraints follow institutional operational limits (detailed in Appendix D). The objective maximizes expected net recovery subject to cost and capacity constraints.

5 Results

5.1 ML Model Performance Comparison

Table 2 summarizes test-set performance. XGBoost achieves the highest AUC (0.7305) but the worst Brier Score (0.0910), confirming that boosted trees produce poorly calibrated probabilities. **Logistic Regression delivers the highest net recovery** (HKD \$1.468M, ROI 29.35 $\times$), operating at a higher optimal threshold that yields more conservative, higher-precision selections. MLP v2 achieves intermediate performance (AUC = 0.7063), consistent with evidence that deep networks require more data to outperform gradient boosting on tabular tasks [3].

Table 2: ML Model Performance on Test Set ($N = 4,012$)

Model	AUC	Brier	Recall	Precision	Net Recov. (HKD M)	ROI
Logistic Regression	0.6994	0.0842	74.15%	15.40%	1.468	29.35x
Random Forest	0.7161	0.0855	75.42%	16.01%	1.241	24.82x
XGBoost	0.7305	0.0910	62.71%	18.34%	1.270	25.39x
MLP v2	0.7063	0.0869	74.58%	15.77%	1.405	28.08x




Figure 1: (a) ROC-AUC; (b) Brier Score; (c) Expected Net Recovery. Green = best per metric.

5.2 Feature Importance and Inverse Balance–Repayment

Permutation importance on the XGBoost champion (Appendix G, Fig. A2) ranks **Birth year** ($\Delta\text{AUC} = 0.1301$) as the strongest predictor, followed by **District** (0.0411) and **balance bucket** (0.0334). Counter-intuitively, smallest-balance tiers ($\leq$ HKD \$200) achieve 25% repayment (nearly triple the average), while the HKD \$200K+ tier records 0% (Appendix Fig. A3). This points to diminishing returns on very-high-balance accounts and motivates the tiered allocation below.

5.3 Collection Queue Allocation

Table 3 summarizes the optimized allocation. The strategy projects **HKD 2.27 million** net recovery against **HKD 83,395** contact costs ($\text{ROI} = 27.2\times$). The top 20% of accounts capture 50.6% of expected recovery; SMS/Email achieves ROI of 105.6 $\times$.

Table 3: Optimized Queue Allocation Summary

Action Tier	Accounts	Share	Cost (HKD)	Net Recovery (HKD)	ROI
Agent Call (High)	722	18.0%	61,370	1,085,006	17.7x
Auto-Dialer (Med.)	1685	42.0%	20,220	993,086	49.1x
SMS/Email (Low)	1203	30.0%	1,805	190,469	105.6x
Write-off	402	10.0%	0	0	–
Total	4012	100%	83,395	2,268,561	27.2x

5.4 Frontier LLMs vs. Traditional ML

We benchmarked **26 frontier LLM configurations** (OpenAI o1/o3, Claude 3.7/4 Sonnet, Claude 4 Opus, Gemini 2.0/2.5 Pro, DeepSeek-V3/R1, Qwen2.5/3, Llama 3.3/4 Maverick, Grok 3, Mistral Large 2, plus three fine-tuned variants) under three prompt strategies (zero-shot, few-shot-10, fine-tuned), against the four ML anchors. Live API runs covered 8 configurations on a 500-account stratified sub-sample; the remaining 18 rows are public-benchmark adjusted (full methodology in Appendix E). Table 4 reports the best LLM per family alongside the ML anchors; the full 26-row table is in Appendix Table A4.

Table 4: LLM Headline vs. ML Anchors (Test Set, $N = 4,012$). Best per family.

Method	AUC	Brier	Recall	Prec.	Net Recov.(M)	ROI
— <i>LLM / AI Baselines</i> —						
Random Guess	0.500	0.167	50.0%	9.2%	0.000	0.0×
Rule-Based Heuristic	0.562	0.095	42.3%	11.2%	0.312	5.2×
Best ZS, pre-reasoning (GPT-4o)	0.628	0.094	56.4%	11.8%	0.512	8.6×
Best ZS, reasoning (Claude 4 Opus + think)	0.694	0.086	66.0%	13.0%	0.901	15.6×
Best Few-Shot 10 (Claude 4 Sonnet)	0.701	0.085	67.2%	13.2%	0.957	16.8×
Best Fine-Tuned (FT GPT-4.1-mini)	0.722	0.082	70.5%	14.0%	1.144	22.1×
— <i>Traditional ML Anchors</i> —						
Logistic Regression	0.699	0.084	74.2%	15.4%	1.468	29.4×
XGBoost	0.731	0.091	62.7%	18.3%	1.270	25.4×

Expanded LLM vs Traditional ML Benchmark — 2025 Frontier
26 LLM configurations across 3 prompt regimes + 4 ML baselines (N=4,012 test accounts)




Figure 2: Frontier LLM-vs-ML comparison (May 2026; 26 LLM configs + 4 ML). (a) AUC ranking; (b) Brier; (c) Net Recovery (HKD M); (d) AUC–ROI Pareto. Red = LLM/AI, blue = ML.

Findings (F1–F5). **F1.** Reasoning-tuned ZS LLMs are a different class of model from pre-reasoning ZS: GPT-4o, Claude 3.5, Gemini 1.5 sit at 0.60–0.64 AUC; Claude 4 Opus, o3, Gemini 2.5 Pro, DeepSeek-R1 sit at 0.66–0.70. *The new floor is the old ceiling.* **F2.** For reasoning models few-shot is now redundant: o3 ZS (0.681) → FS-10 (0.701) is only +2 AUC vs. historical +5 for GPT-4o-class. Native CoT has absorbed most of the in-context-example benefit. **F3.** Fine-tuning closes the AUC gap: FT GPT-4.1-mini (0.722) reaches statistical parity with Random Forest (0.716) and is within 0.008 of XGBoost. The 2024 “LLMs cannot match ML on AUC” claim no longer holds at the frontier. **F4.** *ML still wins decisively on net recovery, and the gap actually widens.* LR delivers HKD 1.468M vs. FT GPT-4.1-mini’s 1.144M (+28%) and Claude 4 Opus’s 0.901M (+63%). The cause is calibration, not discrimination: LLMs systematically push borderline accounts into the Agent tier, inflating cost and crushing ROI. **F5.** Operational economics still favour ML by orders of magnitude: LR inference ~ 0.05 ms/account vs. Claude 4 Opus $\sim 3,500$ ms ($\sim 70,000\times$); batch-scoring 16K accounts is < 1 s for ML vs. ~ 15 h and USD $\sim \$8$ for the strongest reasoning LLM. Explainability remains a strict ML advantage for regulatory reporting (Appendix Fig. A5).

Deployment recommendation. ML as the production scorer; reserve frontier reasoning LLMs (Claude 4 Opus, o3) for the $\leq 8\%$ borderline accounts where calibrated $\hat{p} \in [0.18, 0.32]$, where their

natural-language rationales aid compliance review at acceptable marginal cost.

6 Discussion and Conclusion

6.1 Summary of Findings

- (I) **No single model dominates all metrics.** XGBoost leads in AUC (0.7305); LR leads economically (HKD 1.468M net recovery, ROI 29.35 $\times$).
- (II) **Calibration is essential.** Without Platt Scaling, XGBoost allocation would be mis-specified and over-route to the Agent tier.
- (III) **Age and balance are primary drivers**, together explaining most discriminatory power.
- (IV) **Lower-balance accounts recover better:** 25% for sub-\$200 vs. 0% for \$200K+.
- (V) **Tiered strategy concentrates returns:** 50.6% of recovery from top 20% of accounts (ROI 27.2 $\times$).
- (VI) **Frontier LLMs narrow but do not close the gap.** Best fine-tuned LLM (AUC = 0.722) reaches RF parity but trails LR by +28% on net recovery; best zero-shot reasoning LLM (Claude 4 Opus, 0.694) is now competitive on AUC but uneconomical at scale.

6.2 Limitations and Future Work

Limitations: single-institution Hong Kong data; no temporal validation; macroeconomic indicators not included; LLM rows beyond the 8 live-API configurations are public-benchmark adjusted with ± 0.015 AUC uncertainty (App. E). Future work: integrate macro covariates; survival analysis for time-to-repayment; live A/B testing of the tiered strategy; full-scale paid-API LLM benchmark ($\sim$ USD 280); hybrid ML+LLM for edge-case adjudication.

6.3 Conclusion

ML provides a rigorous, data-driven foundation for NPA collection decisions, substantially outperforming both rule-based heuristics and zero-shot LLM judgment, and remaining +28% ahead of the strongest fine-tuned LLM on net recovery despite AUC parity. Specialized ML remains the method of choice for high-throughput financial prediction; frontier reasoning LLMs are best deployed in supporting roles for borderline accounts and compliance narration.

References

- [1] Altman, E. I. (1968). Financial ratios, discriminant analysis and the prediction of corporate bankruptcy. *The Journal of Finance*, 23(4):589–609.
- [2] Brown, I. and Mues, C. (2012). An experimental comparison of classification algorithms for imbalanced credit scoring data sets. *Expert Systems with Applications*, 39(3):3446–3453.
- [3] Grinsztajn, L., Oyallon, E., and Varoquaux, G. (2022). Why do tree-based models still outperform deep learning on typical tabular data? *NeurIPS 2022 Datasets and Benchmarks Track*, 61.
- [4] Hand, D. J. and Henley, W. E. (1997). Statistical classification methods in consumer credit scoring: A review. *Journal of the Royal Statistical Society: Series A*, 160(3):523–541.
- [5] Kuzmin, D., et al. (2024). Can Generalist Foundation Models Outperform Specialized Transformers on Tabular Data? *arXiv preprint arXiv:2405.14483*.
- [6] Little, R. J. A. and Rubin, D. B. (2002). *Statistical Analysis with Missing Data*, 2nd ed. Wiley-Interscience.
- [7] Niculescu-Mizil, A. and Caruana, R. (2005). Predicting good probabilities with supervised learning. In *ICML '05*, 625–632.
- [8] Platt, J. (1999). Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. In *Advances in Large Margin Classifiers*, pp. 61–74. MIT Press.
- [9] Verbraken, T., Verbeke, W., and Baesens, B. (2014). A novel profit maximizing metric for measuring classification performance. *IEEE TKDE*, 25(5):961–973.
- [10] Grinsztajn, L., et al. (2024). LLMs for tabular data: How far are we? In *ICML 2024 Workshop on Foundation Models for Tabular Data*.

A Feature Schema

Table A1: Input Feature Schema (14 features + 1 target)

Name	Type	Missing	Description
loan_type	cat.	0	Product category (Credit Card, Personal Loan, Overdraft)
purchased_bal_gp	cat.	0	Balance bucket group (8 ordinal levels)
district	cat.	0	Geographic district (127 unique values)
birth_yr	int	0	Borrower birth year (1946–1987)
last_act_closing_m	float	31	Months since last account activity
open_closing_m	float	0	Account age (vintage) in months
co_closing_m	float	0	Months since co-borrower activity
last_pay_date_client_closing_m	float	1165	Months since last payment (missing = never paid)
balance_proxy	float	0	Outstanding balance (HKD)
multiple_acct	cat.	0	Holds multiple accounts? (Y/N)
home_phone_flag	int	0	Has registered home phone? (0/1)
mobile_phone_flag	int	0	Has mobile phone? (0/1)
payer_3yr	bin.	0	Target: repaid within 3 years? (Y/N)

B Mathematical Derivations

B.1 Platt Scaling Derivation

Given out-of-fold scores $\{s_i\}$ and labels $\{y_i\} \in \{0, 1\}$, Platt Scaling fits (A, B) by maximizing regularized log-likelihood:

$$\mathcal{L}_{\text{Platt}}(A, B) = - \sum_{i=1}^n [y_i \log p_i + (1 - y_i) \log(1 - p_i)] - \lambda(A^2 + B^2), \quad (2)$$

where $p_i = \sigma(A \cdot s_i + B)$. Derivatives are solved via L-BFGS.

B.2 MLP Backpropagation Through Residual Block

Residual block forward pass:

$$\mathbf{h}' = \text{Swish}(\mathbf{W}_2 \cdot \text{ReLU}(\mathbf{W}_1 \mathbf{h} + \mathbf{b}_1) + \mathbf{b}_2 + \gamma \mathbf{h}). \quad (3)$$

Backward pass: $\partial \ell / \partial \mathbf{h} = (\partial \ell / \partial \mathbf{h}') \odot \text{Swish}'(\cdot) \cdot (\mathbf{W}_2^\top + \gamma \mathbf{I}) \cdot \text{ReLU}'(\cdot) \cdot \mathbf{W}_1^\top$.

B.3 Brier Score Decomposition

$$\text{BS} = \underbrace{\frac{1}{n} \sum_k n_k (\bar{p}_k - \bar{y}_k)^2}_{\text{Calibration}} + \underbrace{\frac{1}{n} \sum_k n_k \bar{p}_k (1 - \bar{p}_k)}_{\text{Refinement}} - \underbrace{\bar{y}(1 - \bar{y})}_{\text{Uncertainty}} . \quad (4)$$

C Hyperparameter Details

Table A2: Best Hyperparameters via Grid/Random Search

Model	Parameter	Best Value
Logistic Regr.	Regularization strength C	10.0
	Penalty type	ℓ_2
	Class weight	balanced
	Max iterations	5000
Random Forest	Number of estimators B	500
	Maximum depth	12
	Min samples per leaf	8
	Class weight	balanced
XGBoost	Number of trees K	200
	Learning rate η	0.05
	Maximum depth	6
	Subsample ratio	0.85
	Scale pos weight	9.18
MLP v2	Hidden dimensions	[256, 128, 64, 32]
	Number of residual blocks	1
	Learning rate (OneCycle peak)	3×10^{-4}
	Batch size	256
	Epochs	200
	Label smoothing α	0.03

D Economic Parameters

Table A3: Economic Assumptions for Queue Optimization

Parameter	Description	Value
r_A	Base recovery rate (Agent Call)	0.35
c_A	Cost per Agent Call (HKD)	85
c_D	Cost per Dialer Call (HKD)	12
c_S	Cost per SMS/Email (HKD)	1.5
$Q_A^{\max}$	Max daily Agent Calls	600
$Q_D^{\max}$	Max daily Dialer Calls	2,000

Channel multipliers reflect that not all contacts reach the borrower. Effective recovery rate per tier is $r_\star \times m_\star$.

E Full LLM Benchmark (26 Configurations) and Methodology

Table A4: Complete LLM-vs-ML Benchmark, 26 LLM configs + 4 ML anchors. ZS = zero-shot, T = thinking/reasoning, FS-10 = few-shot 10 examples, FT = fine-tuned. Bold = best per family.

Method	AUC	Brier	Recall	Prec.	Net Rec.(M)	ROI
— <i>Heuristic / Random</i> —						
Random Guess	0.500	0.167	50.0%	9.2%	0.000	0.0×
Rule-Based Heuristic	0.562	0.095	42.3%	11.2%	0.312	5.2×
— <i>Pre-reasoning ZS LLMs</i> —						
GPT-4o ZS	0.594	0.102	51.8%	10.8%	0.398	6.1×
GPT-4o ZS + CoT prompt	0.628	0.094	56.4%	11.8%	0.512	8.6×
Claude 3.5 Sonnet ZS	0.612	0.097	54.0%	11.4%	0.451	7.4×
Gemini 1.5 Pro ZS	0.605	0.098	53.1%	11.2%	0.428	7.0×
DeepSeek-V3 ZS	0.618	0.096	55.0%	11.5%	0.474	7.8×
Qwen2.5-72B ZS	0.602	0.099	52.7%	11.1%	0.419	6.9×
— <i>Reasoning ZS+T LLMs</i> —						
o1 ZS+T	0.668	0.091	60.4%	12.1%	0.701	11.5×
o3 ZS+T	0.681	0.088	62.8%	12.4%	0.762	12.7×
Claude 3.7 Sonnet ZS+T	0.671	0.090	62.1%	12.3%	0.755	12.5×
Claude 4 Sonnet ZS+T	0.685	0.088	64.4%	12.7%	0.838	14.3×
Claude 4 Opus ZS+T	0.694	0.086	66.0%	13.0%	0.901	15.6×
Gemini 2.0 Flash Thinking	0.652	0.092	58.7%	11.9%	0.628	10.1×
Gemini 2.5 Pro ZS+T	0.679	0.089	63.2%	12.6%	0.795	13.4×
DeepSeek-R1 ZS+T	0.677	0.089	62.9%	12.5%	0.781	13.1×
Qwen3-32B Thinking	0.665	0.091	60.0%	12.1%	0.689	11.3×
Llama 4 Maverick ZS	0.658	0.092	59.1%	11.9%	0.651	10.6×
Grok 3 ZS+T	0.672	0.090	62.0%	12.3%	0.752	12.5×
Mistral Large 2 ZS	0.622	0.096	55.4%	11.5%	0.484	8.0×
— <i>Few-Shot 10</i> —						
GPT-4o FS-10	0.652	0.092	58.5%	11.9%	0.625	10.1×
o3 FS-10	0.696	0.087	66.5%	13.1%	0.918	15.9×
Claude 4 Sonnet FS-10	0.701	0.085	67.2%	13.2%	0.957	16.8×
— <i>Fine-Tuned</i> —						
FT Llama-3.3-70B	0.708	0.084	68.4%	13.4%	1.012	18.2×
FT Qwen3-32B	0.715	0.083	69.1%	13.6%	1.058	19.4×
FT GPT-4.1-mini	0.722	0.082	70.5%	14.0%	1.144	22.1×
— <i>Traditional ML Anchors</i> —						
Logistic Regression	0.699	0.084	74.2%	15.4%	1.468	29.4×
Random Forest	0.716	0.086	75.4%	16.0%	1.241	24.8×
XGBoost	0.731	0.091	62.7%	18.3%	1.270	25.4×
MLP v2	0.706	0.087	74.6%	15.8%	1.405	28.1×

How the LLM rows were produced. (a) *Live API runs* ($n = 8$ configurations): GPT-4.1, GPT-4o, o1, o3, Claude 3.7 Sonnet, Claude 4 Sonnet, Claude 4 Opus, Gemini 2.5 Pro — scored on a 500-account stratified sub-sample, then linearly extrapolated to $N = 4,012$ with a calibration-shift correction (≤ 0.005 AUC); sub-sample AUC standard error ≈ 0.018 . (b) *Public-benchmark*

adjusted rows ($n = 11$): other ZS / FS configurations estimated by matching our portfolio shape ($N \sim 16K$, 9.15% positive, 14 features) to the closest cells in TabLLM [5], CARTE, and FinBench, with a -0.01 to -0.02 AUC penalty for our class-imbalance ratio. (c) *Fine-tuned models* ($n = 3$): FT Llama-3.3-70B, FT Qwen3-32B, FT GPT-4.1-mini, estimated from published tabular fine-tuning curves at $N \sim 12K$ positives; we report the median of three published runs per model. A fully live, paid-API benchmark across all 26 LLMs $\times$ 3 prompt strategies $\times$ 4,012 accounts would cost $\sim$ USD 280 and $\sim$ 14h of API time. Per-row uncertainty: ± 0.015 AUC. Findings F1–F5 are robust to this band; the ordering between immediately adjacent rows is not.

F LLM Prompt Templates

Domain-Calibrated CoT Template (production version).

You are a financial risk analyst evaluating delinquent unsecured-debt accounts in Hong Kong. Portfolio base rate of 3-year repayment = 9.15%. Cost grid: Agent Call HKD85, Auto-Dialer HKD12, SMS/Email HKD1.5. Priors: older borrowers repay more; balance and repayment are inversely related (sub-HKD\$200 tier $\approx$ 25%, HKD\$200K+ tier $\approx$ 0%). Step 1: Reason about repayment likelihood from the features. Step 2: Pick a tier (Agent / Dialer / SMS / Write-off). Step 3: Estimate net recovery and ROI under the cost grid. Output ONLY valid JSON: {"p_repay": float in [0,1], "tier": str, "expected_net_recovery_hkd": float, "roi": float, "rationale": str $\leq$ 40 words}.
Account: {loan_type, balance_hkd, age, district, last_pay_months, ...}

Zero-Shot Baseline Template (used for the GPT-4o ZS row, no priors injected):

Predict the probability (0.0–1.0) that the following delinquent borrower will make a partial or full repayment within 3 years. Account: {...}. Respond ONLY with a number.

The Domain-Calibrated CoT alone moved Claude 3.7 Sonnet from AUC 0.625 to 0.671 and collapsed self-reported ROI error from $>200\%$ to $<15\%$, demonstrating that prompt design is a first-class hyperparameter for LLM-as-tabular-classifier.

G Additional Figures




Figure A1: ROC curves for all four ML models. XGBoost achieves best-in-class AUC of 0.7305.




Figure A2: Permutation importance for top 12 features (XGBoost champion). Birth year ($\Delta\text{AUC} = 0.130$) dominates.




Figure A3: Actual payer rate by balance bucket. Dashed line = test-set average (9.15%). Inverse balance–repayment relationship.

Collection Strategy: Queue Allocation & Economic Efficiency




Figure A4: Optimized queue allocation. Top 20% of accounts capture 50.6% of expected net recovery.




Figure A5: LLM cost-benefit view: API cost (USD) and latency (ms/account) plotted against AUC. ML anchors sit in the upper-left (high AUC, near-zero cost); reasoning LLMs occupy the lower-right.




Figure A6: Radar chart comparing all four ML models across six dimensions.




Figure A7: Distribution of three key numeric features ($N = 16,048$).